



Microservices Observability Patterns

Paulo Alberto Simões

Developer Evangelist, Data-Driven Microservices with Converged Database

What are Microservices?

If every service has to be updated at the same time it's not loosely coupled

“loosely coupled service-oriented architecture with bounded contexts.”

If you know too much about the surrounding services you don't have a bounded context

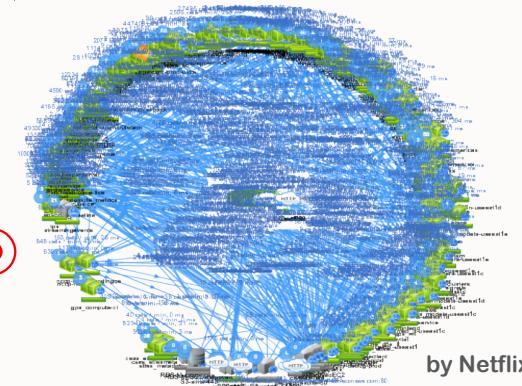
Adrian Cockroft, while at Netflix

Cloud Native Design Principles



Containerized

Open
Source
Software

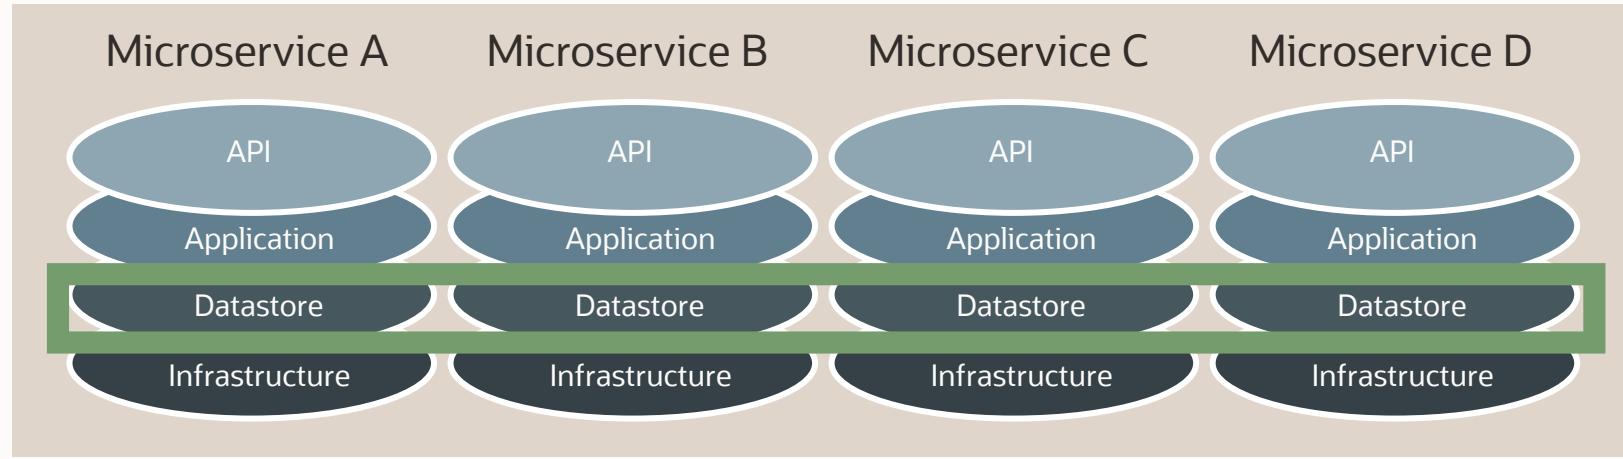


Microservices Oriented



Dynamically orchestrated

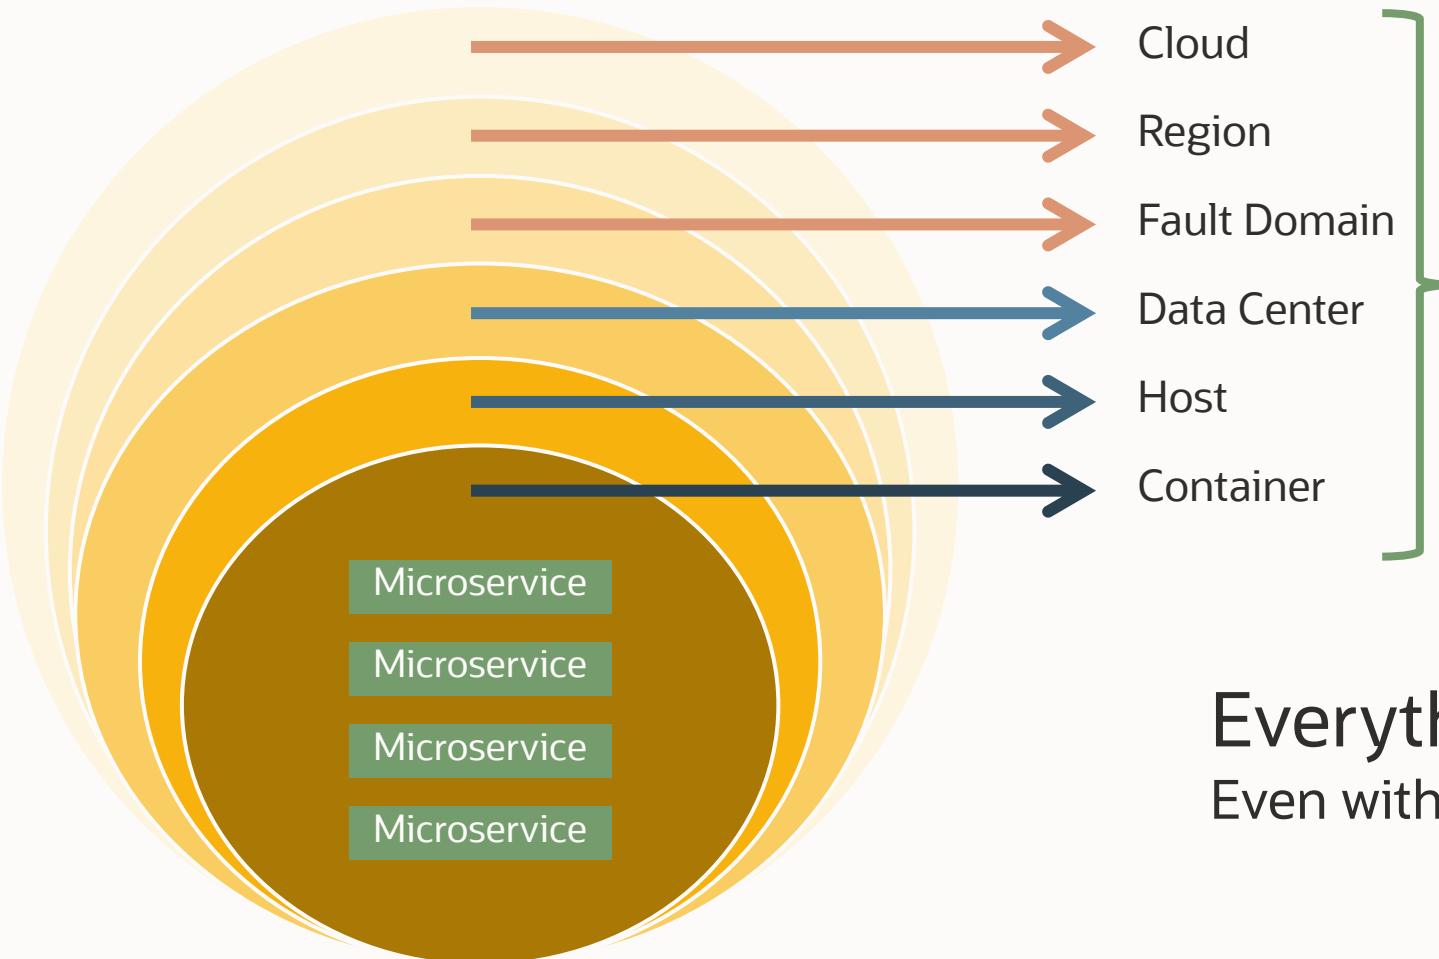
Microservices Forces Move To Distributed Computing



- Distributed computing is a natural consequence of microservices because each microservice has its own datastore
- Sharing datastores across microservices introduces coupling – very bad!
- There will always be latency between microservices
- Latency = eventual consistency

- All data exchange between microservices must be through API layer or messaging – no accessing datastores cross-microservices
- Must implement high-speed messaging between microservices. REST + HTTP probably isn't fast enough
- May end up duplicating data across datastores – e.g. a customer's profile

Distributed Computing forces Everything distributed



Run your applications across multiple data centers, fault domains, regions, etc

Everything is now distributed
Even within the same data center



l.cnfcf.io

This Cloud Native Trail Map is a recommended process for leveraging Open Source, cloud native technologies



CLOUD NATIVE TRAIL MAP

The Cloud Native Landscape l.cnfcf.io has a large number of options. This Cloud Native Trail Map is a recommended process for leveraging open source, cloud native technologies. At each step, you can choose a vendor-supported offering or do it yourself, and everything after step #3 is optional based on your circumstances.

HELP ALONG THE WAY

A. Training and Certification

Consider training offerings from CNCF and then take the exam to become a Certified Kubernetes Administrator or a Certified Kubernetes Application Developer

cncf.io/training

B. Consulting Help

If you want assistance with Kubernetes and the surrounding ecosystem, consider leveraging a Kubernetes Certified Service Provider

cncf.io/kcsp

C. Join CNCF's End User

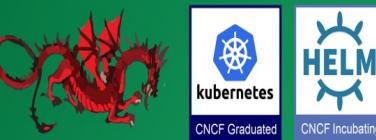
Cloud Native Trail Map

- Commonly done with Docker containers
- Any size application and dependencies (even PDP-11 code running on an emulator) can be containerized
- Over time, you should aspire towards splitting suitable applications and writing future functionality as microservices



3. ORCHESTRATION & APPLICATION DEFINITION

- Kubernetes is the market-leading orchestration solution
- You should select a Certified Kubernetes Distribution, Hosted Platform, or Installer: cncf.io/ck
- Helm Charts help you define, install, and upgrade even the most complex Kubernetes application



5. SERVICE PROXY, DISCOVERY, & MESH

- CoreDNS is a fast and flexible tool that is useful for service discovery
- Envoy and Linkerd each enable service mesh architectures
- They offer health checking, routing, and load balancing



2. CI/CD

- Setup Continuous Integration/Continuous Delivery (CI/CD) so that changes to your source code automatically result in a new container being built, tested, and deployed to staging and eventually, perhaps, to production
- Setup automated rollouts, roll backs and testing

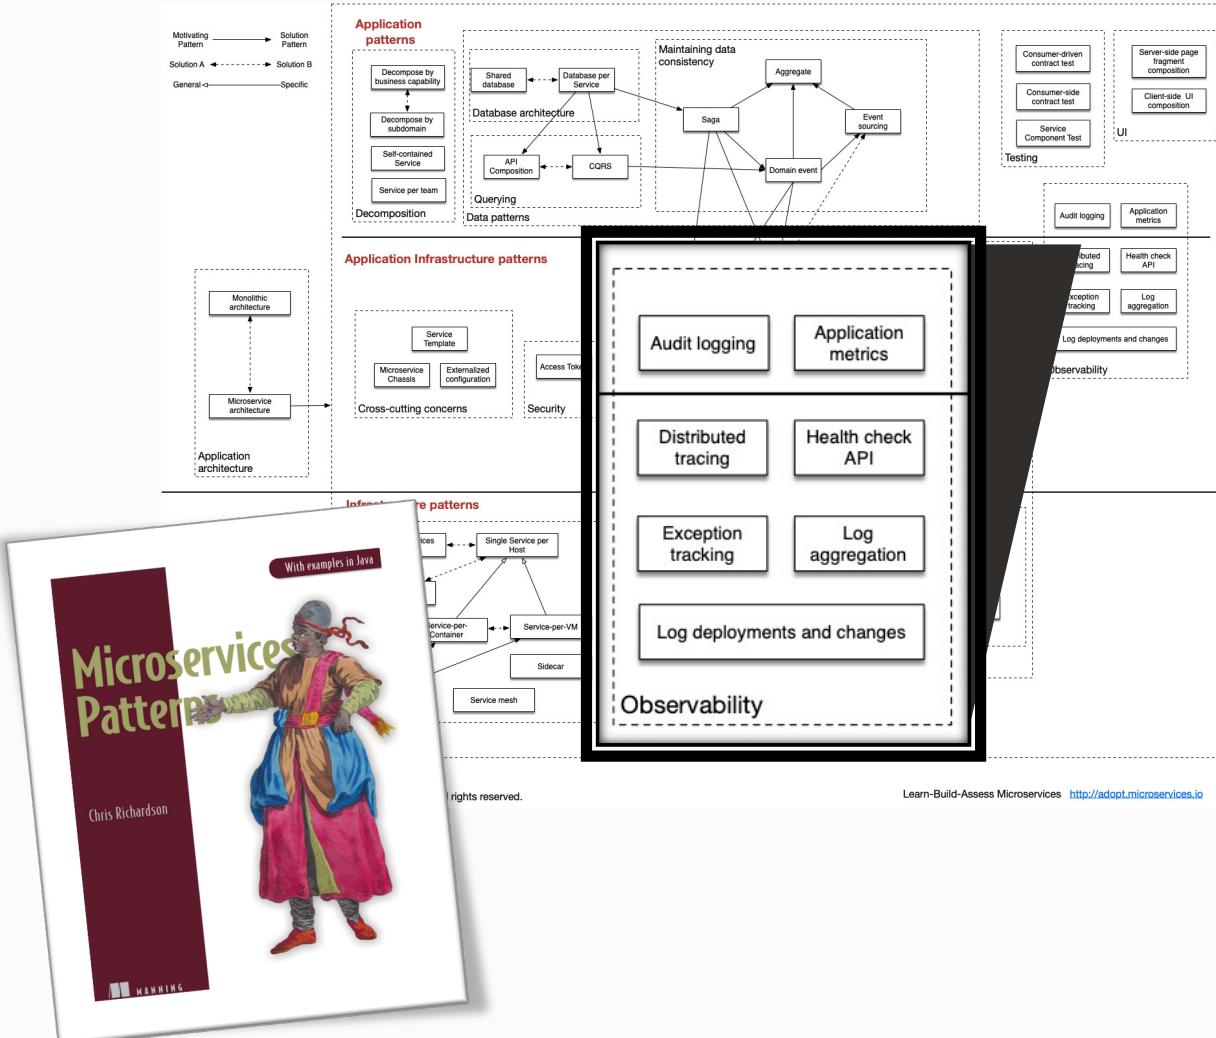
4. OBSERVABILITY & ANALYSIS

- Pick solutions for monitoring, logging and tracing
- Consider CNCF projects Prometheus for monitoring, Fluentd for logging and Jaeger for Tracing
- For tracing, look for an OpenTracing-compatible implementation like Jaeger



Observability Patterns

Microservice Patterns

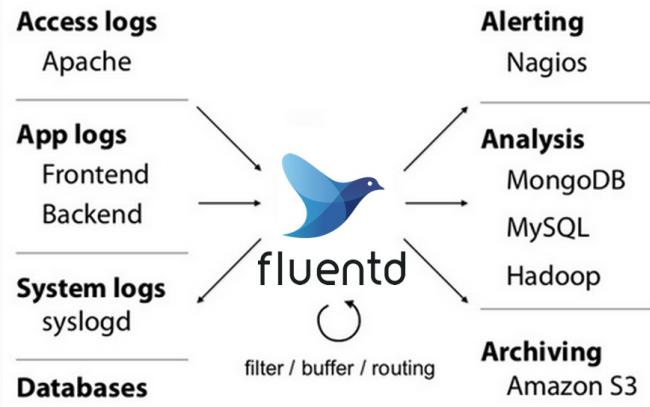
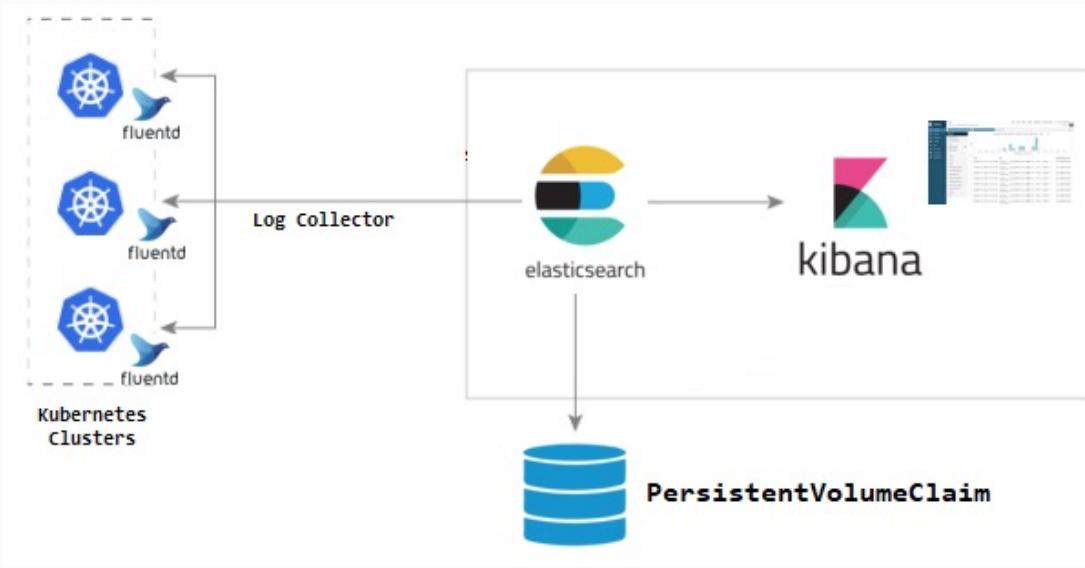


According to microservices architecture and modern systems design there are **7 patterns** to help understand the behavior of an application and troubleshoot problems:

- Log aggregation** - aggregate application logs
- Application metrics** - instrument a service's code to gather statistics about operations
- Audit logging** - record user activity in a database
- Distributed tracing** - instrument services with code that assigns each external request an unique identifier that is passed between services. Record information (e.g. start time, end time) about the work (e.g. service requests) performed when handling the external request in a centralized service
- Exception tracking** - report all exceptions to a centralized exception tracking service that aggregates and tracks exceptions and notifies developers.
- Health check API** - service API (e.g. HTTP endpoint) that returns the health of the service and is intended to be pinged, for example, by a monitoring service
- Log deployments and changes** – CI/CD focus.

Log aggregation

How to understand the behavior of an application and troubleshoot problems?



Oracle Cloud Infrastructure Logging

The screenshot shows the Oracle Cloud Infrastructure Logging interface. At the top, it features the 'fluentd' logo and the 'cloudevents' logo, followed by the 'ORACLE® Cloud Infrastructure' logo. The interface includes search and filter fields: 'FILTER BY FIELD OR TEXT SEARCH' (containing 'type' and sub-options like 'data.apiType') and 'SELECT LOGS TO SEARCH' (containing 'uitest'). Below these are 'Explore' and 'Visualize' tabs. The 'Explore' tab displays a chart titled 'Number of Log Events Per Minute' with data from 14:22 to 17:22 UTC on Sep 28, 2020. The 'Visualize' tab shows log data for the past 3 hours, starting at Mon, Sep 28, 2020, 14:22:15 UTC. Two log entries are visible:

- Mon, Sep 28, 2020, 17:18:57 UTC eventsservice.eventrule.rule... Event failed to deliver to OSS. Exception message: (...)
- Mon, Sep 28, 2020, 17:18:57 UTC eventsservice.eventrule.rule... Event failed to deliver to OSS. Exception message: (...)

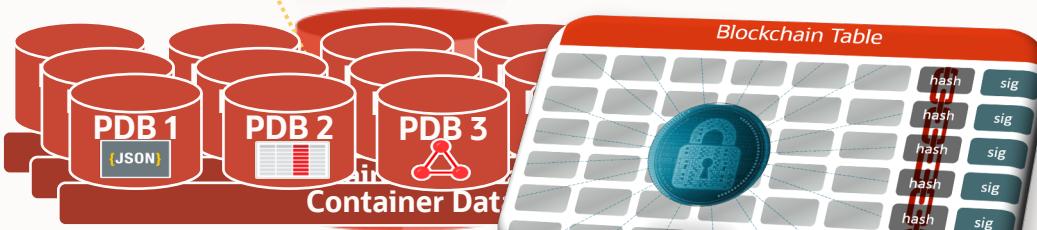


Audit Logging

Keep tracking of actions a user has recently performed (Customer Support, Compliance, etc)



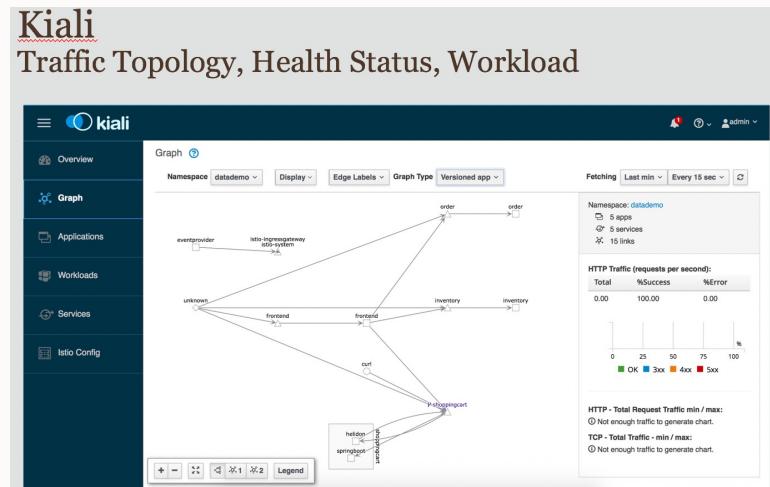
Convergence | Consolidation | Virtualization



Converged Oracle Database

Observability and Distributed Tracing

- Observability: distributed tracing/logging across services
 - Trace, Span, SpanContext, SpanId
- OpenTelemetry: OpenTracing + OpenCensus
- Kiali, Jaeger, and Grafana stop at the edge of the DB
 - Looking into end-to-end distributed tracing



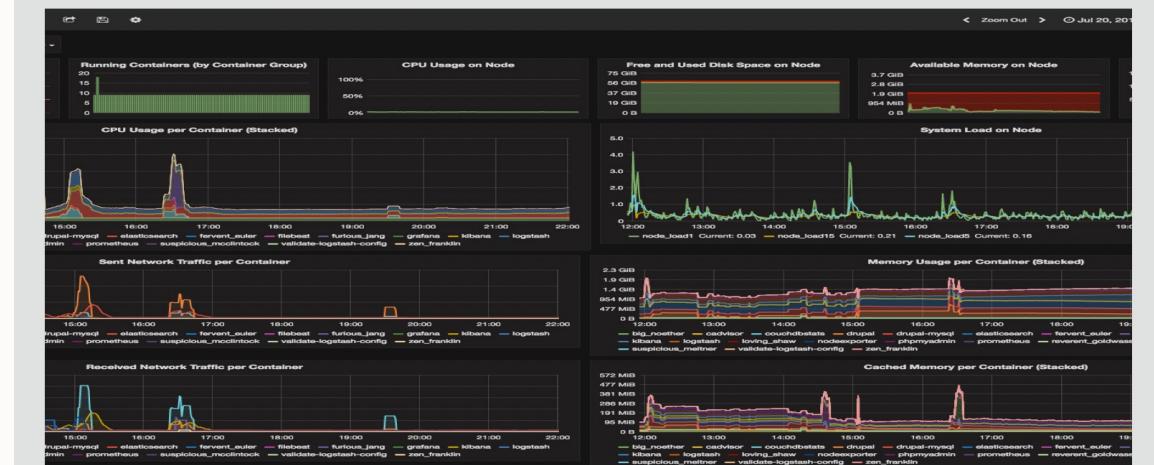
OpenTelemetry - Jaeger

Tracing across full call path of activity
Trace from WebLogic to Helidon MP to Helidon SE (soon to Oracle DB and AQ as well)...

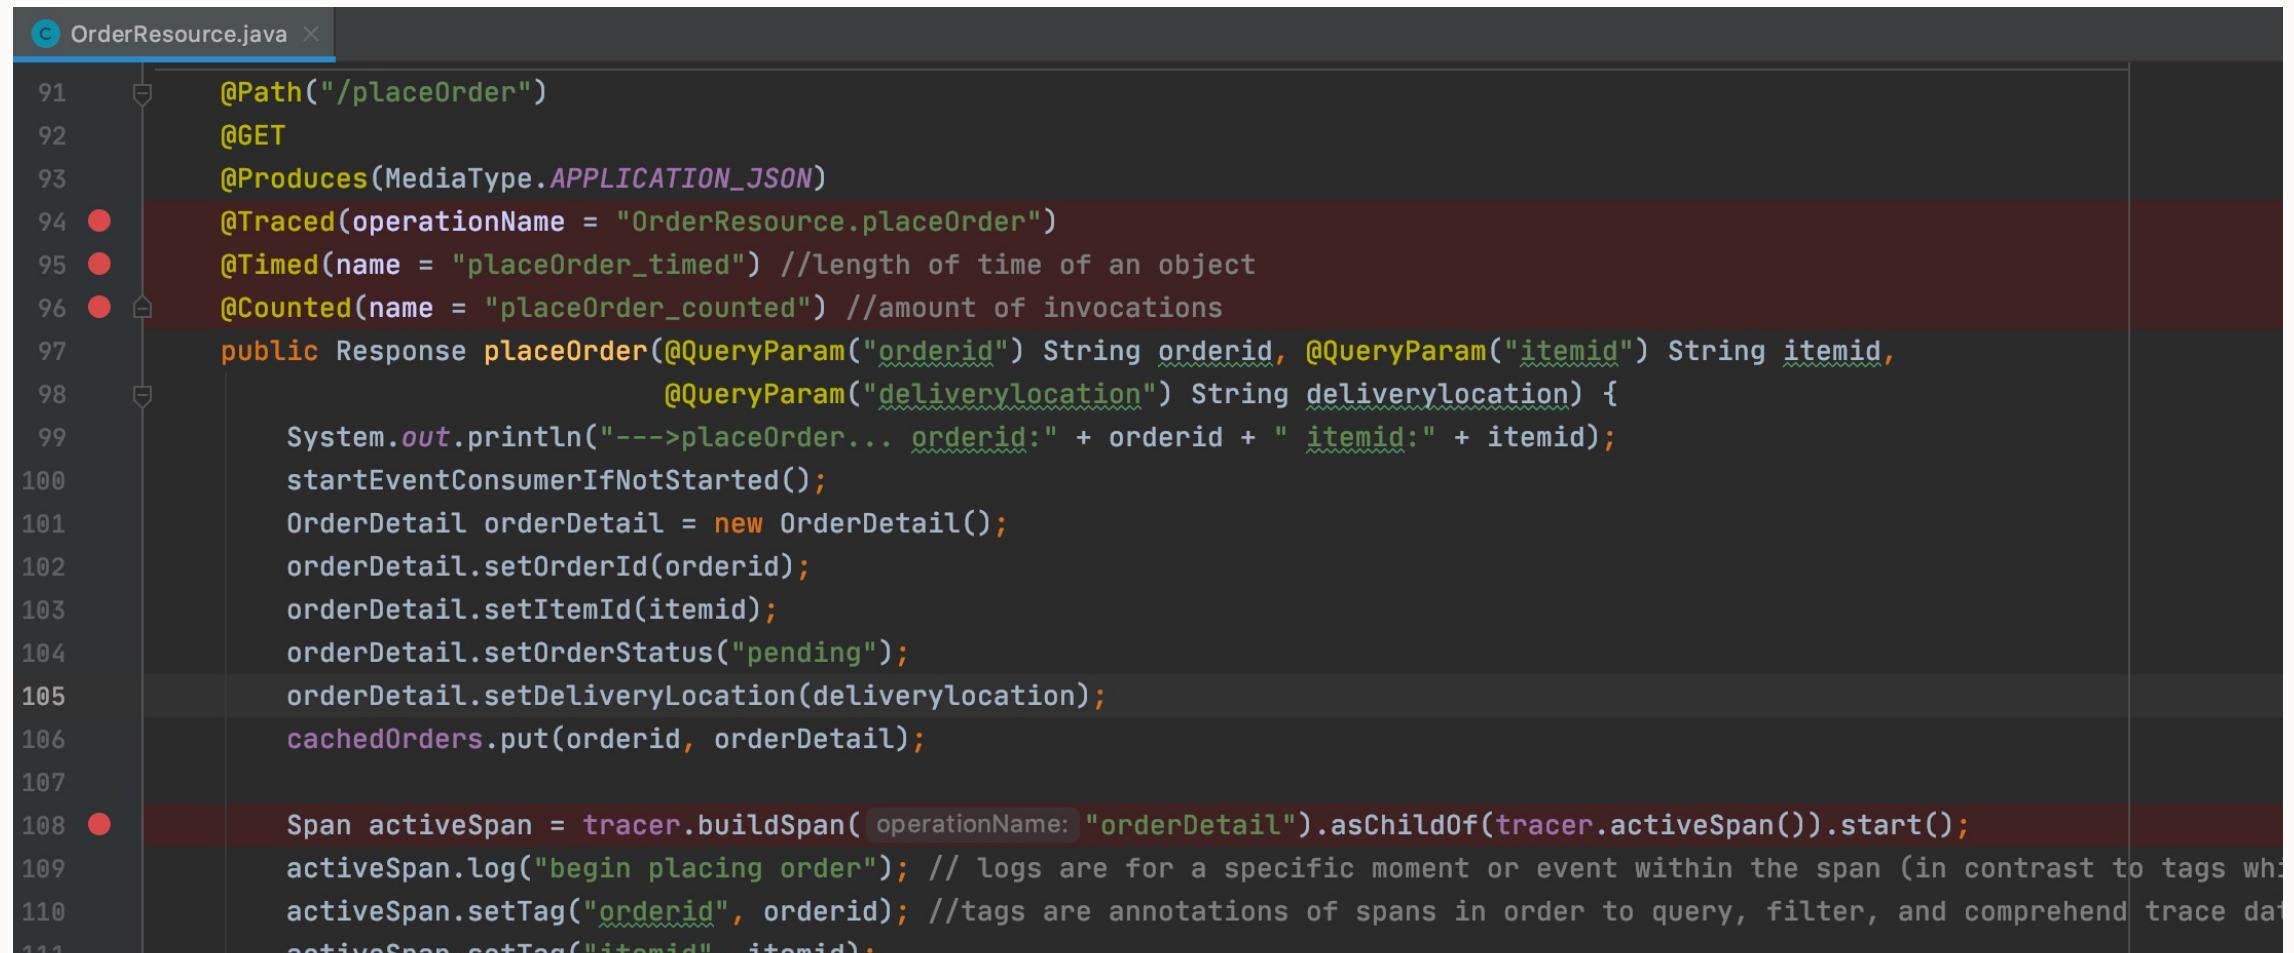


Grafana

Analytics, monitoring, alerting for time-series data



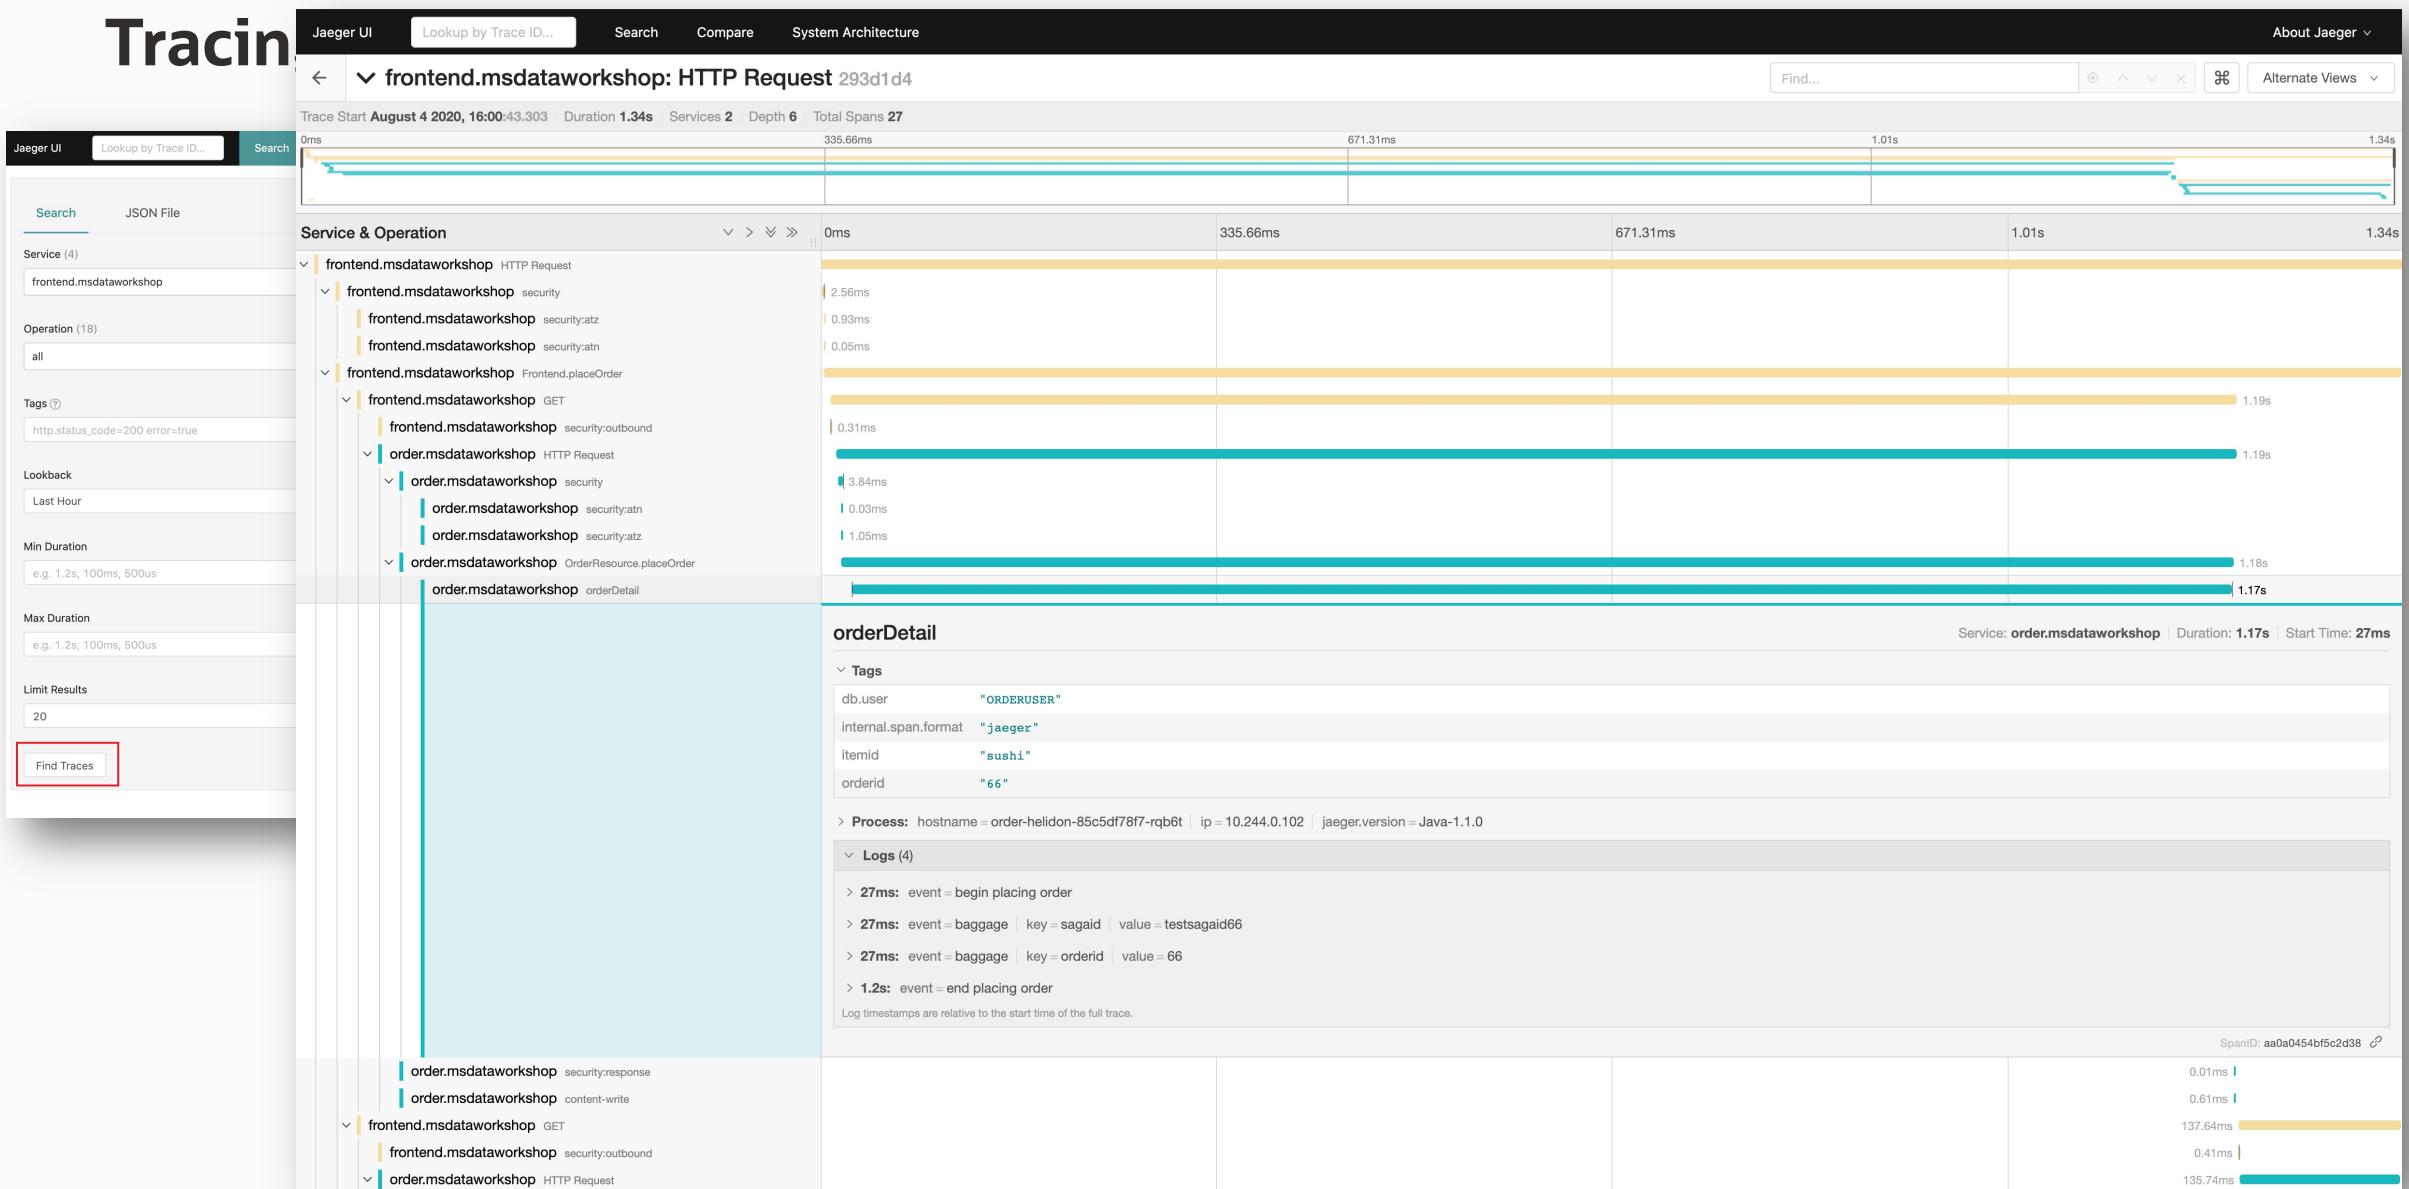
Tracing Using Jaeger



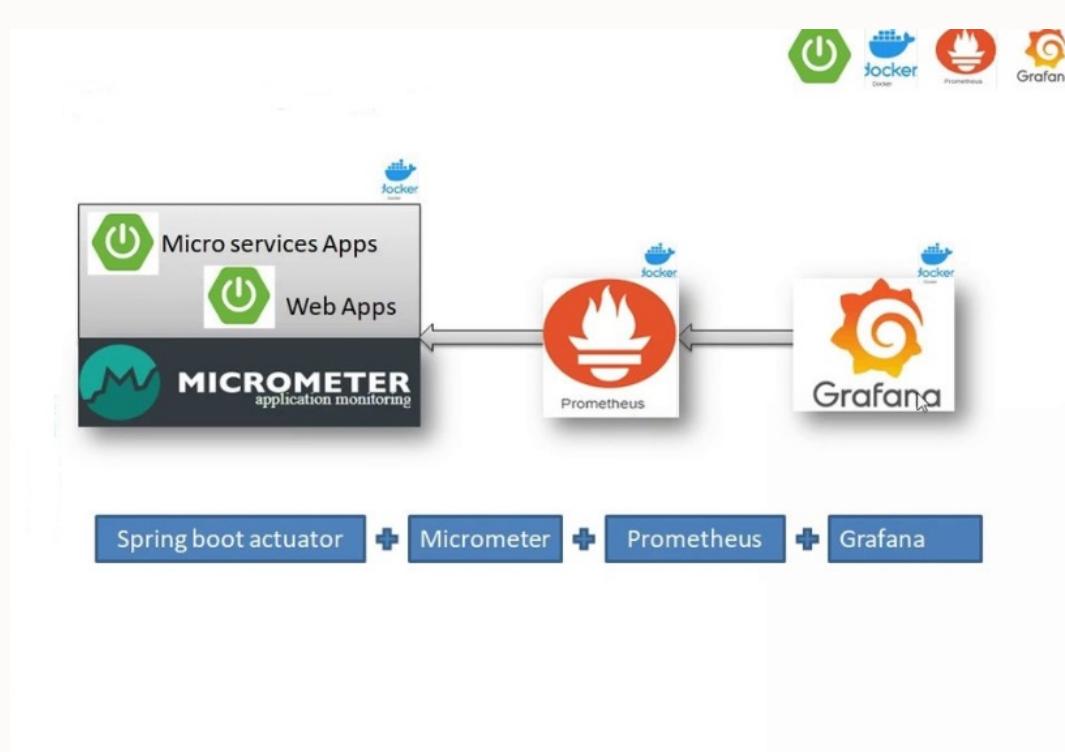
The screenshot shows a Java code editor with the file `OrderResource.java` open. The code defines a REST endpoint for placing an order. Several annotations are used for tracing:

- `@Path("/placeOrder")`
- `@GET`
- `@Produces(MediaType.APPLICATION_JSON)`
- `@Traced(operationName = "OrderResource.placeOrder")`
- `@Timed(name = "placeOrder_timed")` //length of time of an object
- `@Counted(name = "placeOrder_counted")` //amount of invocations
- `public Response placeOrder(@QueryParam("orderid") String orderid, @QueryParam("itemid") String itemid, @QueryParam("deliverylocation") String deliverylocation) {`
- `System.out.println("--->placeOrder... orderid:" + orderid + " itemid:" + itemid);`
- `startEventConsumerIfNotStarted();`
- `OrderDetail orderDetail = new OrderDetail();`
- `orderDetail.setOrderId(orderid);`
- `orderDetail.setItemId(itemid);`
- `orderDetail.setOrderStatus("pending");`
- `orderDetail.setDeliveryLocation(deliverylocation);`
- `cachedOrders.put(orderid, orderDetail);`
- `Span activeSpan = tracer.buildSpan(operationName: "orderDetail").asChildOf(tracer.activeSpan()).start();`
- `activeSpan.log("begin placing order"); // logs are for a specific moment or event within the span (in contrast to tags which are annotations of spans in order to query, filter, and comprehend trace data)`
- `activeSpan.setTag("orderid", orderid); //tags are annotations of spans in order to query, filter, and comprehend trace data`
- `activeSpan.setTag("itemid", itemid);`

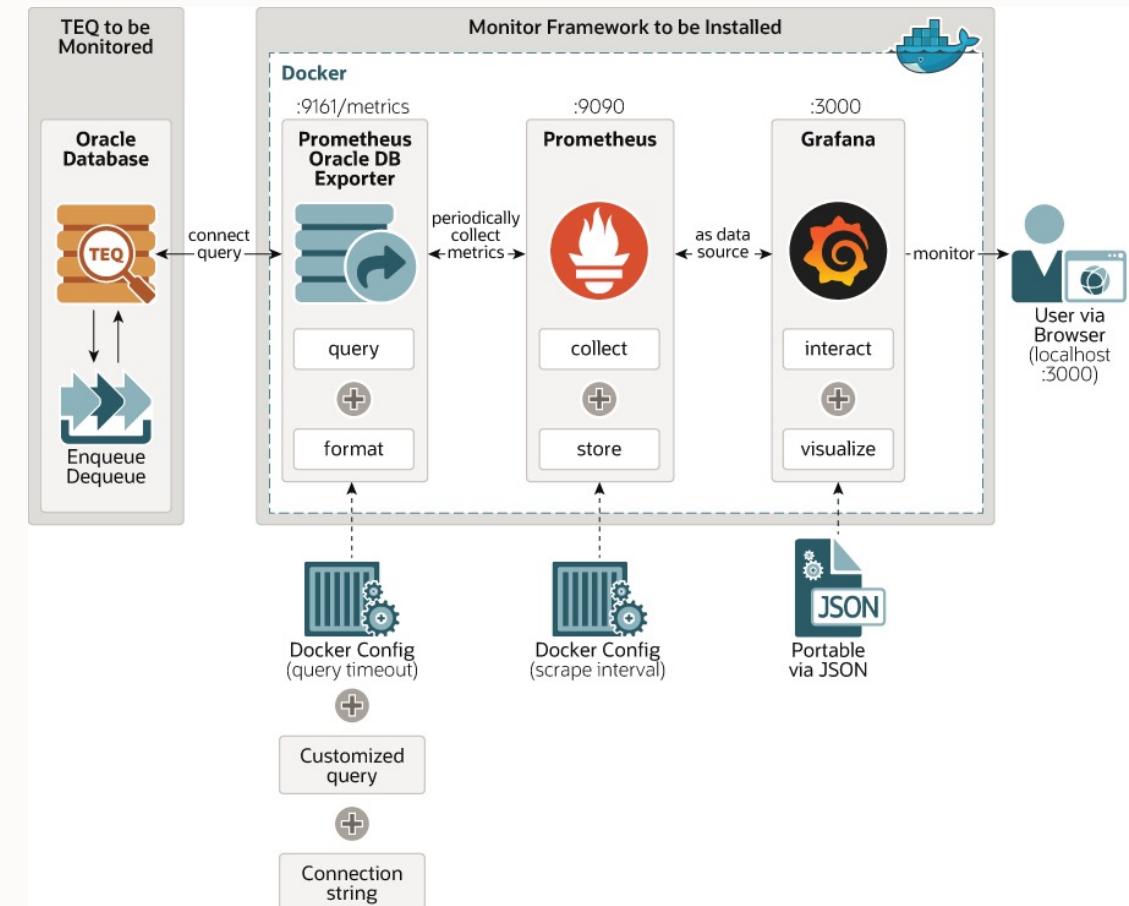
Tracing



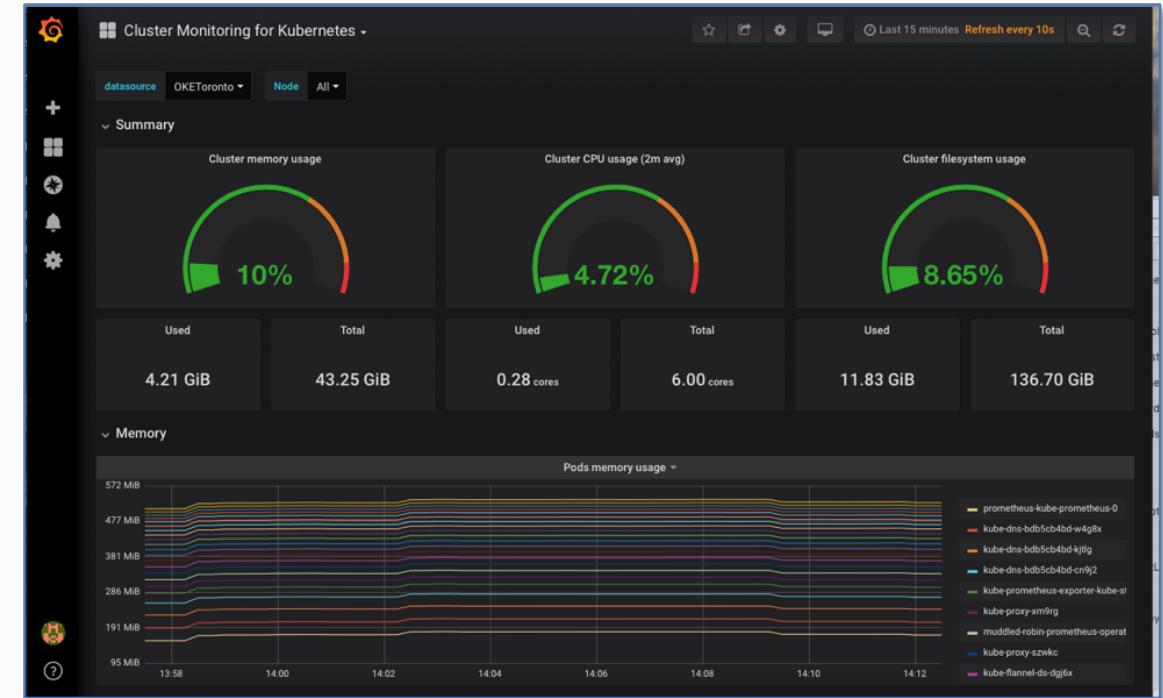
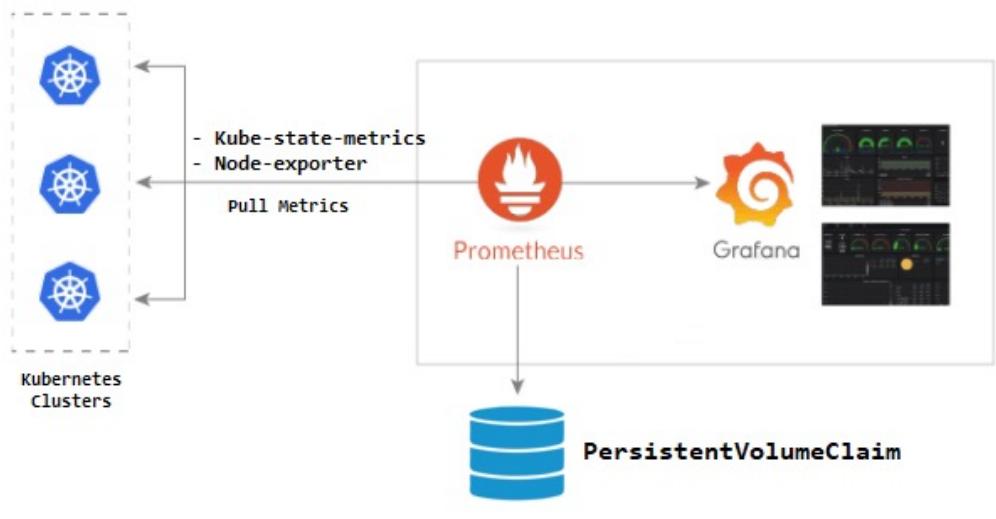
Monitoring Application Microservices & Events (Streams)



Monitoring Oracle DB Transaction Event Queue



Monitoring Infrastructures

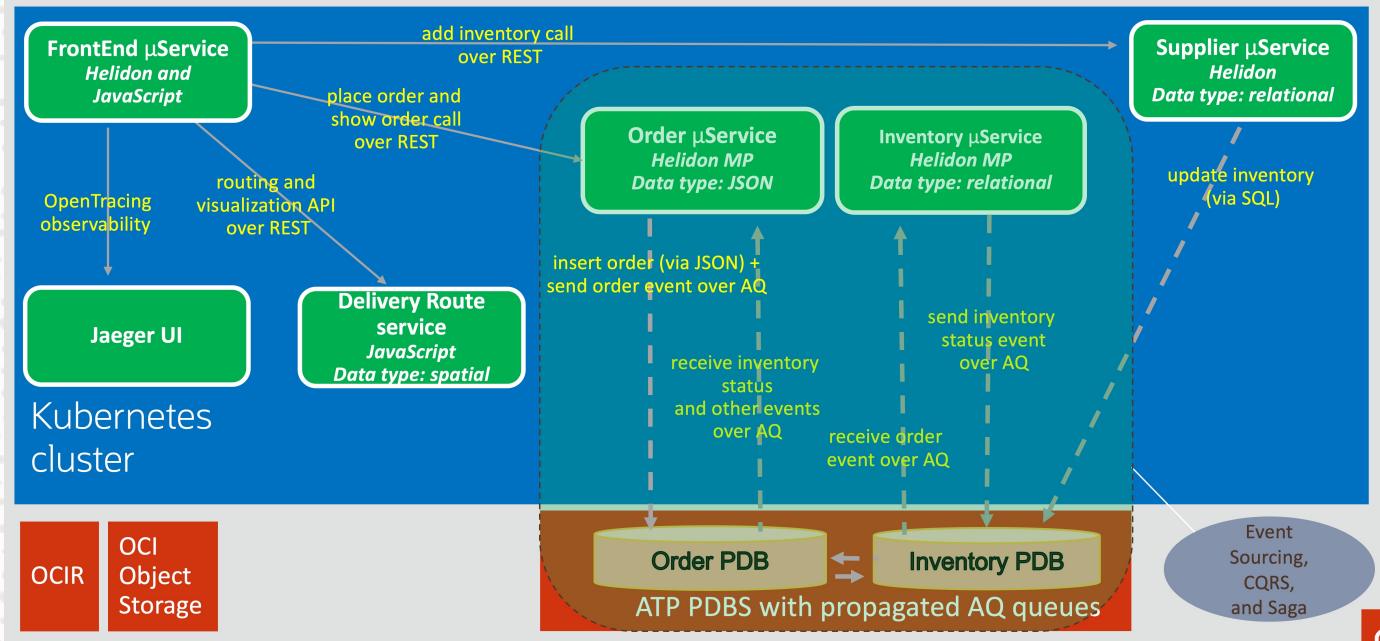


Building Microservices with Oracle Converged Database Workshop

Modernize and simplify your data-driven applications with microservices architecture using Database Cloud and Helidon on OCI.

Workshop length: 2 hours

“GrabDish” Food Order App Architecture



Workshop Outline

- Setup OCI, OKE and ATP
- Build and deploy microservices
- Setup database connection secrets, tables and AQ messaging
- Data-centric microservices walkthrough with Helidon MP
- Polygot Microservices
- Scaling the Application
- Create an APEX App to make sense of the data

<https://bit.ly/simplymicroservices>



Paulo Alberto Simões

Developer Evangelist,
Data-Driven Microservices with Converged Database at Oracle

Mentor, [Oracle for Startups](#) | Ambassador, [Oracle Academy](#)
[Cloud-Native Computing Foundation Ambassador](#)

Stay Connected

LinkedIn, Twitter, Instagram, Clubhouse, twitch.tv, dev.to : **pasimoes**

ORACLE

Our mission is to help people see data in new ways,
discover insights, unlock endless possibilities.